

AutoRAG

PRD Comparison Report

AutoRAG · AutoFounder AI

1

Scoring Summary (out of 10)

Dimension	AutoRAG	AutoFounder AI
Executive Summary	6	9
Problem Statement	7	8
Persona Depth	5	8
Product Goals & Success Metrics	5	8
Scope Definition (In/Out/Deferred)	4	7
Agent Architecture Design	9	9
Feature Specification Clarity	9	8
API & Integration Design	9	6
Data & Storage Design	9	5
Non-Functional Requirements	9	6
AI Safety & Governance	9	6
Monitoring & Observability	5	8
LLMOps / Continuous Learning	4	10
Testing Strategy	9	5
Deployment & DevOps	8	8
UX & UI Requirements	8	7
Timeline & Milestones	7	7
Business Model & Pricing	8	7
Market Analysis (TAM/SAM/SOM)	0	9
GTM Strategy	0	10
Competitive Analysis	0	9
ROI Quantification	0	10
Compliance & Security	7	9
Risks & Assumptions	9	5
Open Questions	8	6

Infrastructure Cost Estimation	6	9
Overall Score	6.9 / 10	7.8 / 10

2 Section-by-Section Coverage

PRD Section	AutoRAG	AutoFounder AI
Executive Summary	Partial (intro only)	Full (vision, differentiators, KPIs)
Problem Statement	✓	✓ + 'Why Now' timing argument
Named Personas	✗ Generic roles	4 named (Sudhanshu, Santosh, Asit, Agency)
Product Goals & KPIs	✗ Missing	TTM, retention D90, CSAT, MoM growth
Scope Definition (In/Out)	✗ Missing	Partial (in-scope + 3 out-of-scope)
System Overview / Architecture	✓	✓
Agent Architecture	3 agents	4 core + 7-pillar expansion
Agent Memory Design	Qdrant, Redis	Redis short-term, Qdrant long-term
Agent Guardrails	Detailed	Partial (PII, injection, content filter)
Feature Breakdown	5 features + specs	5 features (GTM feature best-in-class)
API Design (endpoint list)	14 endpoints	✗ No endpoint list (gRPC only)
External Integrations	Basic (LLMs only)	30+ (GitHub, AWS, Stripe, X, LinkedIn...)
Data & Storage Design	Detailed (Qdrant/MinIO/PG)	✗ Very thin (schema-per-tenant only)
Non-Functional Requirements	Precise (p50/p95/p99)	Partial (incomplete targets)
AI Safety & Governance	Strong (audit log, faithfulness gate)	Partial (bias + HITL, no depth)
Monitoring & Observability	Basic	LangSmith + Prometheus + Grafana
LLMOps / Continuous Learning	✗ Missing	Best-in-class (RLHF, DSPy, Step Fns)
Testing Strategy	Strong (90%+ coverage)	✗ Very thin (3 bullet points)
CI/CD Pipeline	Full gate diagram	Full + CT pipeline (3rd leg)
UX & UI Requirements	20 screens listed	Magic Box, Monaco editor, Launch CC
Timeline & Milestones	Day-by-day	Sprint-based (Week 1-6, 5 phases)
Business Model & Pricing	Rs-based tiers	Rs tiers (■10K / ■50K / Custom)
Market Analysis (TAM/SAM/SOM)	✗ Missing	3 markets with CAGR (\$2.1B→\$12.5B)
GTM Strategy	✗ Missing	Full 8-dimension Market Mapping + CI
Competitive Analysis	✗ Missing	Anticipate moves, gaps, differentiators
ROI Quantification	✗ Missing	Per-pillar ROI (95–99% improvement)

Infrastructure Cost Estimation	Partial	Full 3-tier (100/1K/10K users)
Compliance & Security	GDPR, HIPAA, SOC2	7 frameworks incl. IL AIVITEA + NYC AEDT
Risks & Assumptions	4 categories, 10 named risks	3 brief bullets, no structure
Open Questions	7 prioritised (High/Medium)	20 questions, no owners/dates
7-Pillar Appendix	✗	Full sub-workflow + agent + ROI
CT Pipeline	✗	RLHF + DSPy + AWS Step Functions
Failure Taxonomy (5 types)	✓ Unique	✗
Unified Scoring System	✓ Unique	✗
Pipeline Configurator	✓ Unique	✗
Experiment Tracker + Rollback	✓ Unique	✗
Sign-Off / Reviewer Block	✓ Only one	✗

3

Where AutoRAG Wins Over AutoFounder AI

Dimension	AutoRAG Advantage
Technical Depth	p50/p95/p99 latency targets, per-module coverage targets (90%+), Qdrant indexing strategy, Presidio PII scanner, cost-per-query targets — none of this exists in AutoFounder.
API Design	14 named, typed endpoints with path params and HTTP methods. AutoFounder says 'gRPC internally' with zero endpoint spec.
Data & Storage Design	Full stack: MinIO → PostgreSQL → Qdrant → MLflow — with dimension counts, similarity metrics, retention policy. AutoFounder has 'schema-per-tenant' and nothing else.
Testing Rigor	Module-level coverage targets (90%/85%/80% per agent), schemathesis contract tests, real Qdrant + PostgreSQL in CI. AutoFounder has 3 bullet points.
Risks & Assumptions	4 categories, 10 named risks with probability/impact. AutoFounder has 3 bullets with no structure.
Failure Taxonomy	5-type failure classification is genuinely original IP — AutoFounder has no equivalent.
Unified Scoring System	Single score across multiple quality dimensions — product differentiator with no parallel in AutoFounder.
Experiment Tracker	MLflow-backed versioned pipeline comparison — AutoFounder doesn't touch experiment management.
Open Questions Quality	7 prioritised questions (High/Medium) tied to unresolved MVP decisions. AutoFounder's 20 questions are all strategic with no owners.
Reviewer Sign-Off	Formal technical review block — professional standard missing from AutoFounder.

4 Where AutoFounder AI Wins Over AutoRAG

Dimension	AutoFounder AI Advantage
GTM Engineering	Entire Section 20 with 8-dimension Market Mapping + CI framework — AutoRAG has zero GTM content.
Market Analysis	3-segment TAM/SAM/SOM with CAGR. AutoRAG has no market sizing.
ROI Tables	Per-pillar ROI: 95–99% time/cost reductions with before/after numbers. AutoRAG has no ROI proof points.
LLMOps / CT Pipeline	RLHF loop + DSPy prompt optimisation + model routing + drift detection + nightly A/B evaluation — a full continuous training engine.
7-Pillar Appendix	Each pillar is a mini-PRD: sub-workflows, agents, integrations, ROI — extremely investor-ready.
Named Personas	Sudhanshu / Santosh / Asit / Agency with specific pain stories. AutoRAG says 'ML Engineers' without faces.
Infrastructure Cost	100/1K/10K user cost tiers with per-user economics. AutoRAG gives a monthly infra cap only.
Monitoring & Observability	LangSmith + Prometheus + Grafana + code quality drift detection — AutoRAG's monitoring section is thin.
Compliance Breadth	7 frameworks including IL AIVITEA and NYC AEDT law — shows global market awareness.
Executive Summary	Vision, differentiators, business impact targets, and KPIs all in Section 1 — sets the stage before technical depth.

5 Top 5 Specific Improvements for AutoRAG PRD

#	What to Add	How AutoFounder Does It
1	ROI Proof Table AutoRAG cuts RAG evaluation from 3 days → 15 minutes	AutoFounder's per-pillar ROI tables (before/after time, cost, quality %)
2	Market Sizing RAG evaluation market TAM/SAM/SOM with CAGR	AutoFounder Section 19.1 (3 rows: market, 2026 value, 2030 value, CAGR)
3	LLMOps / CT Pipeline Nightly eval, drift detection, prompt optimisation loop	AutoFounder Section 23.1 (CI/CD + CT as 3rd pillar with RLHF + Step Functions)
4	Named Personas 'Neha the MLOps Engineer at Series B, spends 3 days manually QA-ing RAG'	AutoFounder Section 3 persona table (name, role, key need)
5	GTM Strategy Where does AutoRAG get its first 50 users?	AutoFounder Section 20 (Product Hunt, HN, GitHub, AI engineer Discord, Euron)

6

One-Line Verdicts

AutoRAG 6.9 / 10

The strongest technical PRD spec of the two — the failure taxonomy, unified scoring, and testing rigor are genuinely production-grade — but it reads like it was written for engineers, not stakeholders. Zero market, zero GTM, zero ROI proof.

AutoFounder AI 7.8 / 10

The strongest business and go-to-market document of the two — the 7-pillar ROI appendix and LLMOps CT pipeline are best-in-class — but the technical foundations (API design, data storage, testing, risks) are thin enough that an engineering team couldn't build from this alone.

9–10 Excellent

7–8 Good

5–6 Fair

0–4 Needs Work